


CLIP Architecture

Contrastive Language-Image Pre-training



Model: CLIP-ViT-B/32

Parameters: ~151M

Tokenized text → Embedding

Context: text length up to 77 tokens